

Assignment 3: K-means Clustering Algorithm with Python

Section: ZBB

Group 3:

- Eunice Chua
- Tareq Haboukh
- Fides Audrielle Urgel
- Liezel Anne Viray

Clustering

The purpose of this assignment is to use Python to learn how to perform K-means clustering in Python, and find the optimal value of K.

Instructions

Using Python, you are to complete the following questions. Please submit your answers (CODE USED AND OUTPUT) as PDF files.

Please answer following questions:

1. Find your preferred dataset from Kaggle which is appropriate for an unsupervised learning problem.

<https://www.kaggle.com/>

Dataset used for this assignment is from [IBM HR Analytics Attrition](#).

1. Explore the dataset and provide information about that. You are free to use any preprocessing tools that you want. You can explain the problem and the purpose of the dataset. Visualizing is the best approach to exploring the dataset.

This dataset contains the details of 1,470 IBM employees which includes seniority level, satisfaction rates as well as their attrition data. This dataset could be used by Human Resource department to analyze the performance and satisfaction of employees. It could also help in identifying contributing factors that lead to employees leaving the organization. Kmeans clustering will aim to identify groups of employees with similar attributes. This classification model could be used for further analysis and model building to better predict employee attrition rate.

1. Perform K-means clustering algorithm on your dataset with a range of values for K to choose the optimal value with Elbow method.

- Import K-means from sklearn.cluster.
- Apply K-means on the dataset and get y_pred.
- Calculate the WSS. You can write your own function or get the inertia_ attribute from the fitted model.
- Calculate the silhouette score by using: silhouette_score(X(actual), y_kmeans(predicted)).
- Plot the values of K vs WSS.
- Plot the output clusters with the optimal K.
- Plot the centers of the clusters on the previous plot and show the centroids with a larger size.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
```

Data Importing

```
In [2]: # Importing the IBM HR Analytics Attrition dataset
Attr_df = pd.read_csv("Attrition.csv")
Attr_df.head()
```

Out[2]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Education | EducationField | EmployeeCount | EmployeeNumber | ... | Rela |
|---|-----|-----------|-------------------|-----------|------------------------|------------------|-----------|----------------|---------------|----------------|-----|------|
| 0 | 41 | Yes | Travel_Rarely | 1102 | Sales | 1 | 2 | Life Sciences | 1 | 1 | ... | |
| 1 | 49 | No | Travel_Frequently | 279 | Research & Development | 8 | 1 | Life Sciences | 1 | 2 | ... | |
| 2 | 37 | Yes | Travel_Rarely | 1373 | Research & Development | 2 | 2 | Other | 1 | 4 | ... | |
| 3 | 33 | No | Travel_Frequently | 1392 | Research & Development | 3 | 4 | Life Sciences | 1 | 5 | ... | |
| 4 | 27 | No | Travel_Rarely | 591 | Research & Development | 2 | 1 | Medical | 1 | 7 | ... | |

5 rows × 35 columns



Data Preprocessing

In [3]:

Attr_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                  1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                       1470 non-null   int64
17  MaritalStatus                        1470 non-null   object
18  MonthlyIncome                        1470 non-null   int64
19  MonthlyRate                          1470 non-null   int64
20  NumCompaniesWorked                   1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                    1470 non-null   int64
24  PerformanceRating                    1470 non-null   int64
25  RelationshipSatisfaction              1470 non-null   int64
26  StandardHours                        1470 non-null   int64
27  StockOptionLevel                     1470 non-null   int64
28  TotalWorkingYears                    1470 non-null   int64
29  TrainingTimesLastYear                1470 non-null   int64
30  WorkLifeBalance                      1470 non-null   int64
31  YearsAtCompany                       1470 non-null   int64
32  YearsInCurrentRole                   1470 non-null   int64
33  YearsSinceLastPromotion               1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

In [4]:

```
# drop unnecessary variables: EmployeeCount, EmployeeNumber, StandardHours, Over18
Attr_df.drop(['EmployeeCount', 'EmployeeNumber', 'StandardHours', 'Over18'], axis=1, inplace=True)

Attr_df.shape
```

Out[4]: (1470, 31)

In [5]:

```
# check if there are any Missing data points
Attr_df.isna().sum()
```

Out[5]:

| | |
|--------------------------|---|
| Age | 0 |
| Attrition | 0 |
| BusinessTravel | 0 |
| DailyRate | 0 |
| Department | 0 |
| DistanceFromHome | 0 |
| Education | 0 |
| EducationField | 0 |
| EnvironmentSatisfaction | 0 |
| Gender | 0 |
| HourlyRate | 0 |
| JobInvolvement | 0 |
| JobLevel | 0 |
| JobRole | 0 |
| JobSatisfaction | 0 |
| MaritalStatus | 0 |
| MonthlyIncome | 0 |
| MonthlyRate | 0 |
| NumCompaniesWorked | 0 |
| Overtime | 0 |
| PercentSalaryHike | 0 |
| PerformanceRating | 0 |
| RelationshipSatisfaction | 0 |
| StockOptionLevel | 0 |
| TotalWorkingYears | 0 |
| TrainingTimesLastYear | 0 |
| WorkLifeBalance | 0 |
| YearsAtCompany | 0 |
| YearsInCurrentRole | 0 |
| YearsSinceLastPromotion | 0 |
| YearsWithCurrManager | 0 |

dtype: int64

```
In [6]: # Transform categorical into numerical variable using Label encoder
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()

categorical = Attr_df.select_dtypes(include='object').columns

print('Transforming the following {} Categorical variables into numerical:'.format(len(categorical)))

for i in categorical:
    print(i)
    Attr_df[i]=label_encoder.fit_transform(Attr_df[i]).astype('int64')

Attr_df[categorical].head().style.hide_index()
```

Transforming the following 8 Categorical variables into numerical:
Attrition
BusinessTravel
Department
EducationField
Gender
JobRole
MaritalStatus
OverTime

Out[6]:

| Attrition | BusinessTravel | Department | EducationField | Gender | JobRole | MaritalStatus | OverTime |
|-----------|----------------|------------|----------------|--------|---------|---------------|----------|
| 1 | 2 | 2 | 1 | 0 | 7 | 2 | 1 |
| 0 | 1 | 1 | 1 | 1 | 6 | 1 | 0 |
| 1 | 2 | 1 | 4 | 1 | 2 | 2 | 1 |
| 0 | 1 | 1 | 1 | 0 | 6 | 1 | 1 |
| 0 | 2 | 1 | 3 | 1 | 2 | 1 | 0 |

```
In [7]: Attr_df.dtypes
```

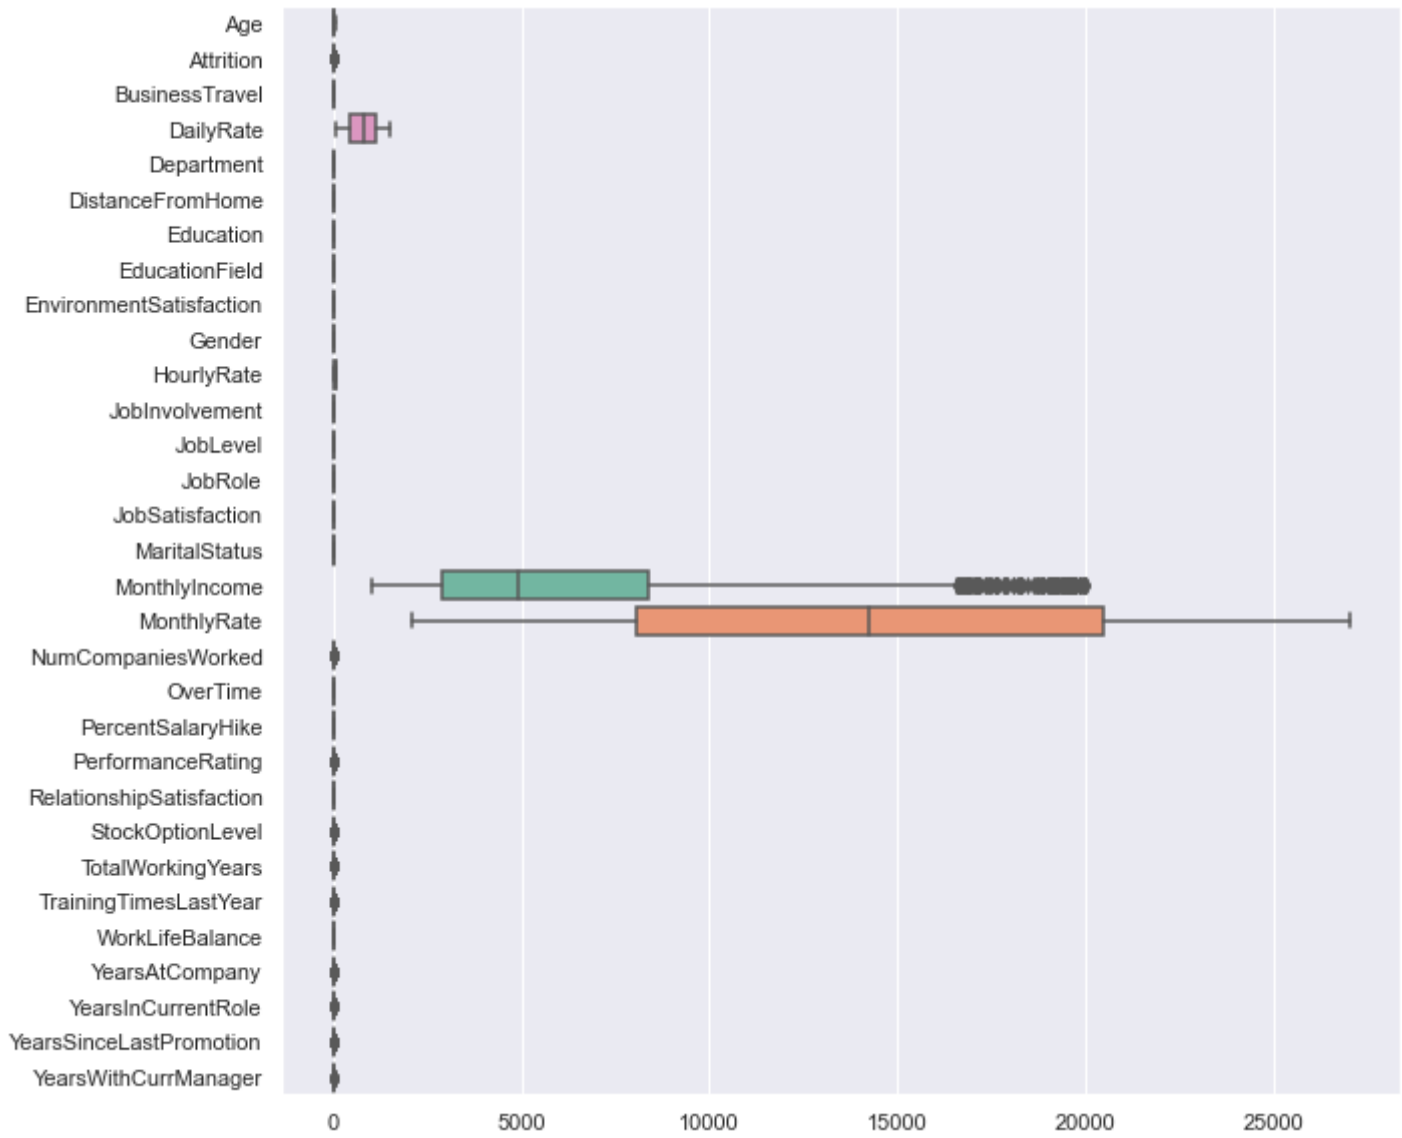
```
Out[7]: Age                int64
Attrition                int64
BusinessTravel           int64
DailyRate                int64
Department               int64
DistanceFromHome         int64
Education                int64
EducationField           int64
EnvironmentSatisfaction  int64
Gender                   int64
HourlyRate               int64
JobInvolvement           int64
JobLevel                 int64
JobRole                  int64
JobSatisfaction           int64
MaritalStatus            int64
MonthlyIncome            int64
MonthlyRate              int64
NumCompaniesWorked       int64
OverTime                 int64
PercentSalaryHike        int64
PerformanceRating        int64
RelationshipSatisfaction  int64
StockOptionLevel         int64
TotalWorkingYears        int64
TrainingTimesLastYear    int64
WorkLifeBalance          int64
YearsAtCompany           int64
YearsInCurrentRole        int64
YearsSinceLastPromotion  int64
YearsWithCurrManager     int64
dtype: object
```

```
In [8]: # check unique values per column
Attr_df.nunique()
```

```
Out[8]: Age                43
Attrition                2
BusinessTravel           3
DailyRate                886
Department               3
DistanceFromHome         29
Education                5
EducationField           6
EnvironmentSatisfaction  4
Gender                   2
HourlyRate               71
JobInvolvement           4
JobLevel                 5
JobRole                  9
JobSatisfaction           4
MaritalStatus            3
MonthlyIncome            1349
MonthlyRate              1427
NumCompaniesWorked       10
OverTime                 2
PercentSalaryHike        15
PerformanceRating        2
RelationshipSatisfaction  4
StockOptionLevel         4
TotalWorkingYears        40
TrainingTimesLastYear    7
WorkLifeBalance          4
YearsAtCompany           37
YearsInCurrentRole        19
YearsSinceLastPromotion  16
YearsWithCurrManager     18
dtype: int64
```

```
In [9]: # boxplot
df_1= Attr_df.loc[:, Attr_df.columns]

sns.set(rc={'figure.figsize':(10,10)})
sns.boxplot(data=df_1, orient="h", palette="Set2");
```



Data Scaling

```
In [10]: # Scaling
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
scaler.fit(df_1)
scaled_df1 = pd.DataFrame(scaler.transform(df_1), columns= df_1.columns)
scaled_df1.head()
```

Out[10]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Education | EducationField | EnvironmentSatisfaction | Gender | ... | I |
|---|-----------|-----------|----------------|-----------|------------|------------------|-----------|----------------|-------------------------|-----------|-----|---|
| 0 | 0.446350 | 2.280906 | 0.590048 | 0.742527 | 1.401512 | -1.010909 | -0.891688 | -0.937414 | -0.660531 | -1.224745 | ... | |
| 1 | 1.322365 | -0.438422 | -0.913194 | -1.297775 | -0.493817 | -0.147150 | -1.868426 | -0.937414 | 0.254625 | 0.816497 | ... | |
| 2 | 0.008343 | 2.280906 | 0.590048 | 1.414363 | -0.493817 | -0.887515 | -0.891688 | 1.316673 | 1.169781 | 0.816497 | ... | |
| 3 | -0.429664 | -0.438422 | -0.913194 | 1.461466 | -0.493817 | -0.764121 | 1.061787 | -0.937414 | 1.169781 | -1.224745 | ... | |
| 4 | -1.086676 | -0.438422 | 0.590048 | -0.524295 | -0.493817 | -0.887515 | -1.868426 | 0.565311 | -1.575686 | 0.816497 | ... | |

5 rows × 31 columns

Feature Selection

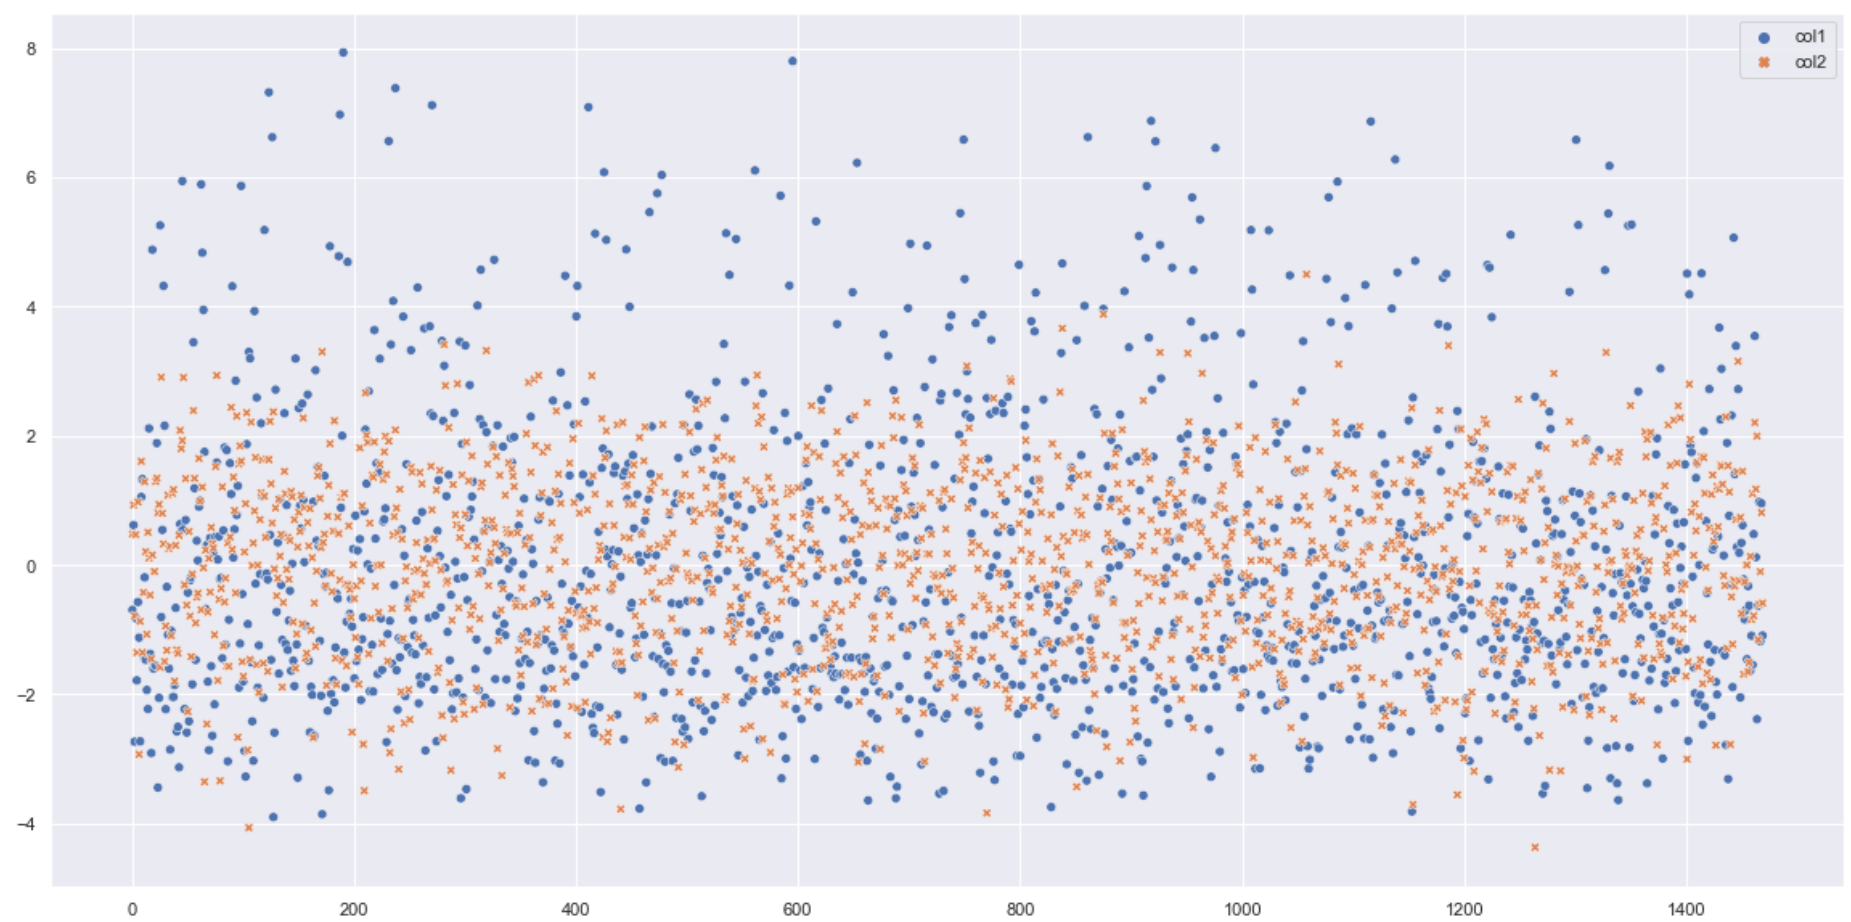
```
In [11]: #PCA: reduce features to 2 dimension only
from sklearn.decomposition import PCA

pca_attr = PCA(n_components=2)
pca_attr.fit(scaled_df1)
PCA_df1 = pd.DataFrame(pca_attr.transform(scaled_df1), columns=(["col1", "col2"]))
PCA_df1.describe().T
```

Out[11]:

| | count | mean | std | min | 25% | 50% | 75% | max |
|------|--------|---------------|----------|-----------|-----------|-----------|----------|----------|
| col1 | 1470.0 | 7.099385e-18 | 2.176639 | -3.901737 | -1.586664 | -0.400616 | 1.097578 | 7.927667 |
| col2 | 1470.0 | -5.452932e-17 | 1.382002 | -4.368912 | -0.981629 | 0.034401 | 1.030095 | 4.495623 |

```
In [12]: # original data set but reduced dimension scatter plot
sns.set(rc={'figure.figsize':(20,10)})
sns.scatterplot(data=PCA_df1);
```



Data Clustering

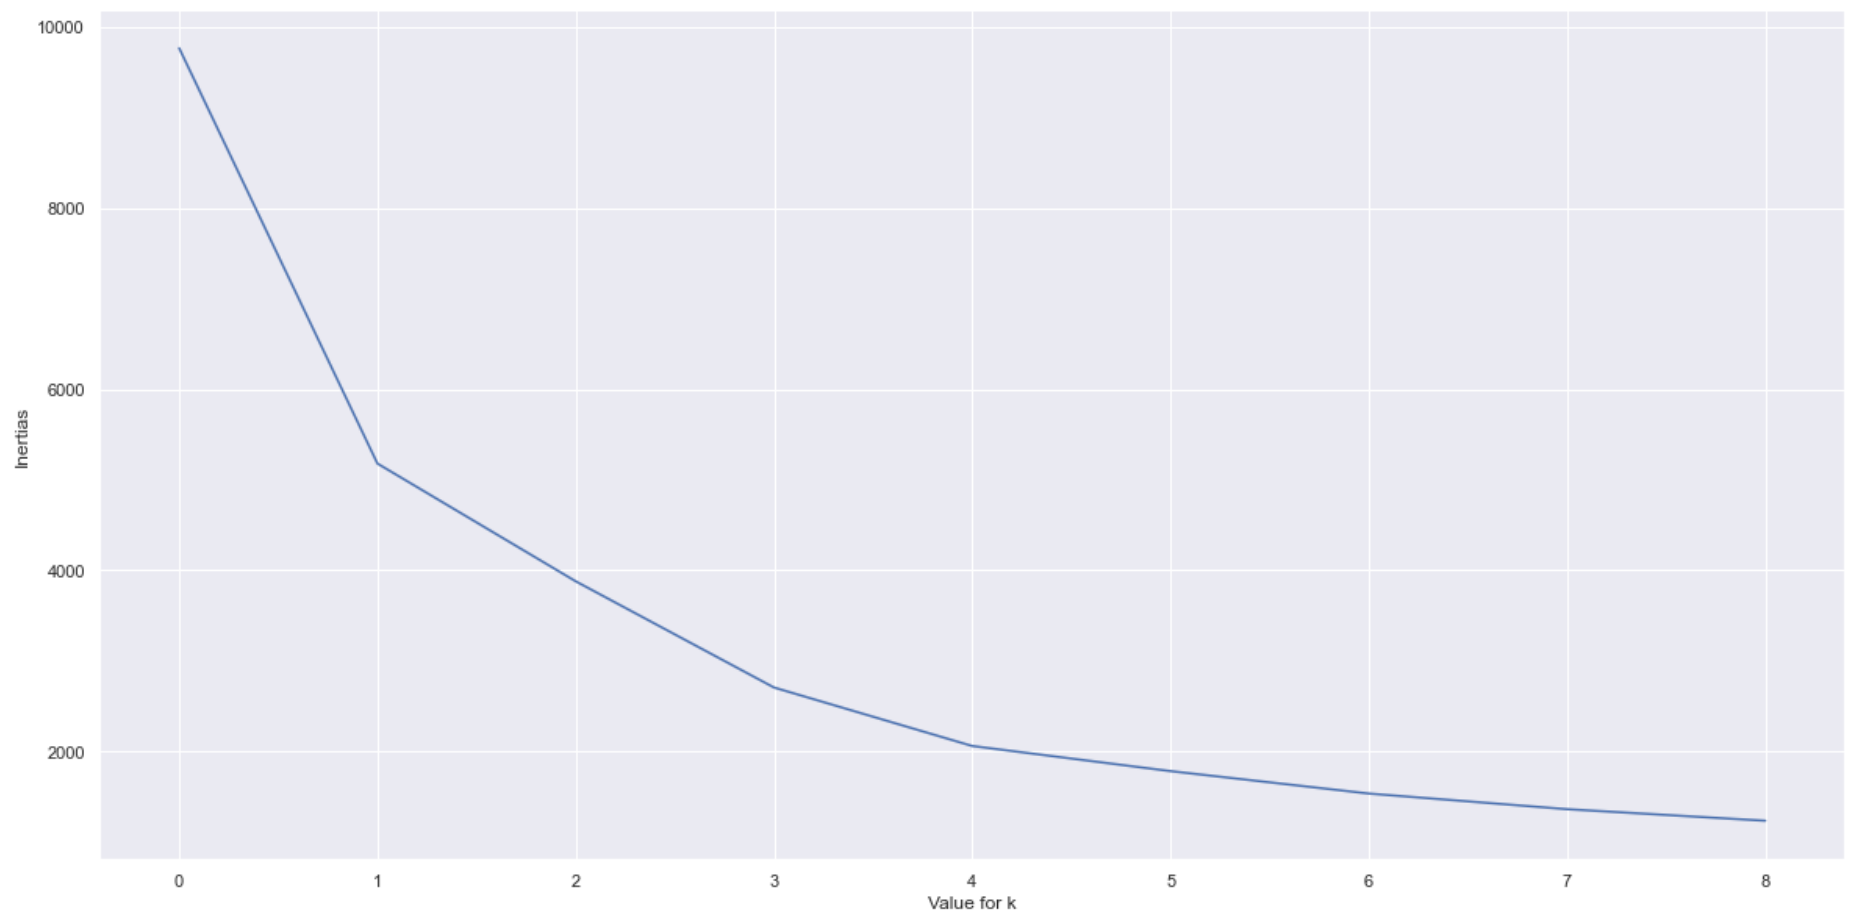
```
In [13]: inertia = {}

for i in range(1, 10):
    kmeans = KMeans(n_clusters=i, max_iter=1000)
    kmeans.fit(PCA_df1)
    inertia[i] = kmeans.inertia_

for k, v in inertia.items():
    print(str(k), ': ', str(round(v,2)))
```

```
# Plot for each K value
plt.subplots()
plt.plot(list(inertia.values()))
plt.xlabel("Value for k")
plt.ylabel("Inertias")
plt.show()
```

```
1 : 9765.46
2 : 5178.18
3 : 3877.85
4 : 2705.24
5 : 2059.29
6 : 1781.74
7 : 1534.08
8 : 1361.41
9 : 1233.85
```



```
In [14]: kmeans = KMeans(n_clusters=4 , n_init=15, random_state=10)
kmeans.fit(PCA_df1)
```

```
Out[14]: KMeans(n_clusters=4, n_init=15, random_state=10)
```

```
In [15]: centroids = kmeans.cluster_centers_  
labels = kmeans.labels_  
pred_clusters = kmeans.predict(PCA_df1)
```

```
In [16]: # inertia score  
print('\nInertia in the K-means clustering',round(kmeans.inertia_,2))  
  
#silhoutte score  
print('\nSilhoutte score for K-means clustering',round(silhouette_score(PCA_df1,pred_clusters),4))  
  
Inertia in the K-means clustering 2705.24  
  
Silhoutte score for K-means clustering 0.3835
```

```
In [17]: # dataframe with Label attached for kmeans clustering  
Attr_dfkmean = PCA_df1.copy()  
Attr_dfkmean['label'] = labels  
Attr_dfkmean.head()
```

```
Out[17]:
```

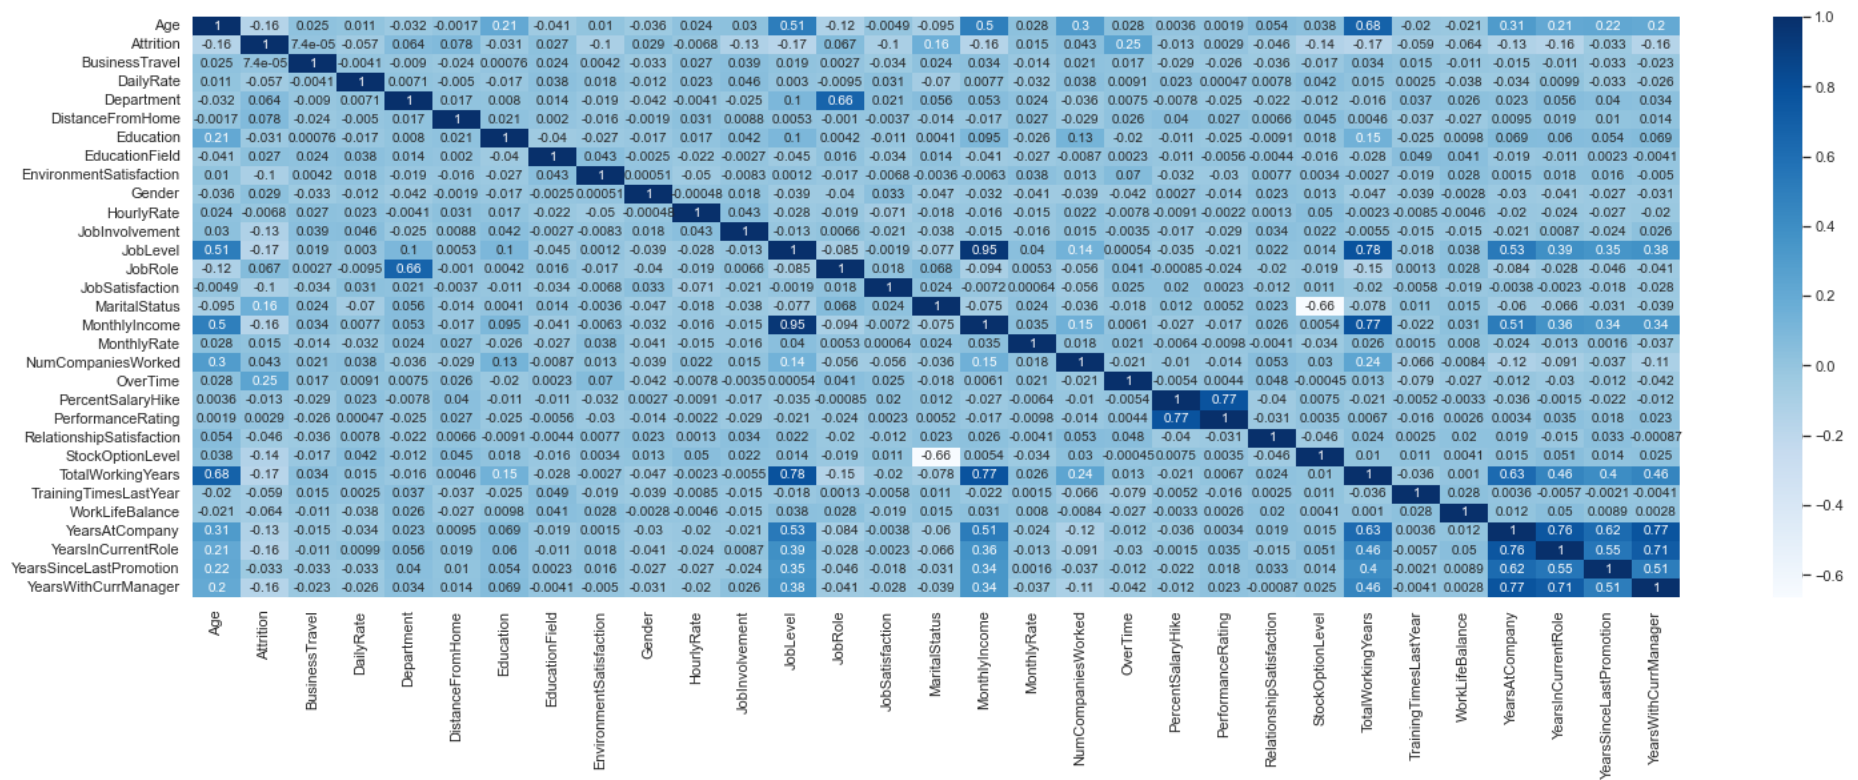
| | col1 | col2 | label |
|---|-----------|-----------|-------|
| 0 | -0.696076 | 0.478191 | 2 |
| 1 | 0.616857 | 0.929435 | 1 |
| 2 | -2.735275 | -0.788440 | 2 |
| 3 | -0.822928 | 0.469346 | 2 |
| 4 | -1.786719 | -1.355992 | 2 |

```
In [18]: plt.figure(figsize=(15,10))  
plt.scatter(Attr_dfkmean[Attr_dfkmean.columns[0]],Attr_dfkmean[Attr_dfkmean.columns[1]],c=kmeans.labels_ , cmap=plt.cm.Set1)  
plt.scatter(centroids[:, 0], centroids[:, 1], c='black', s=200, alpha=0.5)  
plt.show()
```



KMeans without PCA

```
In [19]: %matplotlib inline  
plt.figure(figsize=(25,8))  
sns.heatmap(scaled_df1.corr(),annot=True,cmap='Blues');
```

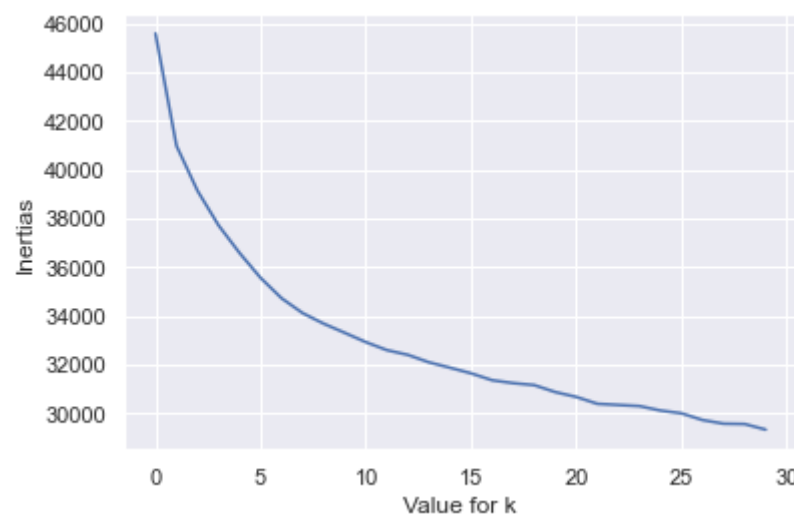
```
In [20]: inertia = {}

for i in range(1, 31):
    kmeans = KMeans(n_clusters=i, max_iter=1000)
    kmeans.fit(scaled_df1)
    inertia[i] = kmeans.inertia_

for k, v in inertia.items():
    print(str(k), ': ', str(round(v,2)))

# Plot for each K value
plt.subplots()
plt.plot(list(inertia.values()))
plt.xlabel("Value for k")
plt.ylabel("Inertias")
plt.show()
```

```
1 : 45570.0
2 : 40965.41
3 : 39119.87
4 : 37702.47
5 : 36567.85
6 : 35545.92
7 : 34709.66
8 : 34102.06
9 : 33669.48
10 : 33294.31
11 : 32913.86
12 : 32585.57
13 : 32391.24
14 : 32084.9
15 : 31864.9
16 : 31634.6
17 : 31351.17
18 : 31232.45
19 : 31151.88
20 : 30865.84
21 : 30666.56
22 : 30385.04
23 : 30337.12
24 : 30291.81
25 : 30107.95
26 : 29994.41
27 : 29719.03
28 : 29569.76
29 : 29554.96
30 : 29325.69
```



```
In [21]: kmeans = KMeans(n_clusters=11 , n_init=15, random_state=10)
kmeans.fit(scaled_df1)
```


Out[21]: KMeans(n_clusters=11, n_init=15, random_state=10)

```
In [22]: centroids = kmeans.cluster_centers_  
pred_clusters = kmeans.predict(scaled_df1)  
labels = kmeans.labels_
```

```
In [23]: # inertia score  
print('\nInertia in the K-means clustering',round(kmeans.inertia_,2))  
  
# silhouette score  
print('\nSilhouette score for K-means clustering',round(silhouette_score(scaled_df1,pred_clusters),4))
```

Inertia in the K-means clustering 32879.26

Silhouette score for K-means clustering 0.0466

```
In [24]: scaled_df1['cluster_kmeans'] = labels  
scaled_df1.head()
```

Out[24]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Education | EducationField | EnvironmentSatisfaction | Gender | ... | I |
|---|-----------|-----------|----------------|-----------|------------|------------------|-----------|----------------|-------------------------|-----------|-----|---|
| 0 | 0.446350 | 2.280906 | 0.590048 | 0.742527 | 1.401512 | -1.010909 | -0.891688 | -0.937414 | -0.660531 | -1.224745 | ... | |
| 1 | 1.322365 | -0.438422 | -0.913194 | -1.297775 | -0.493817 | -0.147150 | -1.868426 | -0.937414 | 0.254625 | 0.816497 | ... | |
| 2 | 0.008343 | 2.280906 | 0.590048 | 1.414363 | -0.493817 | -0.887515 | -0.891688 | 1.316673 | 1.169781 | 0.816497 | ... | |
| 3 | -0.429664 | -0.438422 | -0.913194 | 1.461466 | -0.493817 | -0.764121 | 1.061787 | -0.937414 | 1.169781 | -1.224745 | ... | |
| 4 | -1.086676 | -0.438422 | 0.590048 | -0.524295 | -0.493817 | -0.887515 | -1.868426 | 0.565311 | -1.575686 | 0.816497 | ... | |

5 rows × 32 columns

```
In [25]: fig, axs = plt.subplots(1,3, figsize = (20,5))  
  
plt.subplot(1,3,1)  
sns.scatterplot(data=scaled_df1, x="MonthlyIncome", y="JobLevel", hue="cluster_kmeans");  
  
plt.subplot(1,3,2)  
sns.scatterplot(data=scaled_df1, x="PercentSalaryHike", y="PerformanceRating", hue="cluster_kmeans");  
  
plt.subplot(1,3,3)  
sns.scatterplot(data=scaled_df1, x="TotalWorkingYears", y="JobLevel", hue="cluster_kmeans");
```

